

VLSI System Design HW1

[112061533] [潘金生]

(1.3.1)

Layer (type:depth-idx)	Output Shape	Param #
Net	[4, 10]	—
└─Sequential: 1-1	[4, 6, 28, 28]	—
└─Conv2d: 2-1	[4, 6, 28, 28]	150
└─ReLU: 2-2	[4, 6, 28, 28]	—
└─Sequential: 1-2	[4, 6, 14, 14]	—
└─MaxPool2d: 2-3	[4, 6, 14, 14]	—
└─Sequential: 1-3	[4, 16, 10, 10]	—
└─Conv2d: 2-4	[4, 16, 10, 10]	2,400
└─ReLU: 2-5	[4, 16, 10, 10]	—
└─Sequential: 1-4	[4, 16, 5, 5]	—
└─MaxPool2d: 2-6	[4, 16, 5, 5]	—
└─Sequential: 1-5	[4, 120, 1, 1]	—
└─Conv2d: 2-7	[4, 120, 1, 1]	48,000
└─ReLU: 2-8	[4, 120, 1, 1]	—
└─Sequential: 1-6	[4, 84]	—
└─Linear: 2-9	[4, 84]	10,080
└─ReLU: 2-10	[4, 84]	—
└─Sequential: 1-7	[4, 10]	—
└─Linear: 2-11	[4, 10]	840
Total params: 61,470		
Trainable params: 61,470		
Non-trainable params: 0		
Total mult-adds (M): 1.67		
Input size (MB): 0.02		
Forward/backward pass size (MB): 0.21		
Params size (MB): 0.25		
Estimated Total Size (MB): 0.47		

(1.3.2)

Type	Input activation size (Channel, width, height)	Output activation size (Channel, width, height)	Activation function
Convolution 1	(1, 32, 32) = 6144	(6, 28, 28) = 4704	ReLU
Maxpooling 2	(6, 28, 28) = 4704	(6, 14, 14) = 1176	
Convolution 3	(6, 14, 14) = 1176	(16, 10, 10) = 1600	ReLU
Maxpooling 4	(16, 10, 10) = 1600	(16, 5, 5) = 400	
Convolution 5	(16, 5, 5) = 400	(120, 1, 1) = 120	ReLU
Fully-connected layer 6	(120, 1, 1) = 120	(84)	ReLU
Fully-connected Output	84	10	

(1.3.3)

本次作業架構把 Lenet-5 中的 subsampling layer 換成 maxpooling layer，並且把所有 activation function 換成 ReLu function，且 param 的部分有些差異，因為我們沒有考慮到 bias 的情況。

(1.3.4)

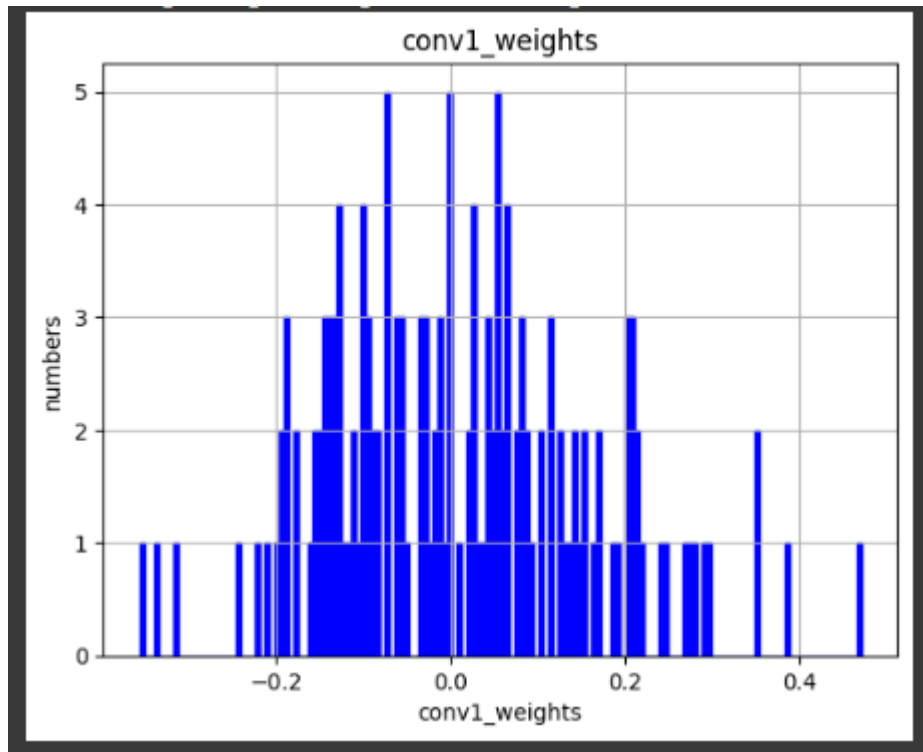
要將 convolution layer 換成 fully connected layer，只要注意向量維度即可，將特徵圖攤平成向量(16*5*5,120)，但這樣可能會遇到一些問題，由於 convolution layer 擁有局部特徵萃取的能力，換成 fully connected layer 會失去局部特徵萃取，再者，這樣會造成參數量大幅增加，提高運算量，造成較大的計算成本，由於參數量增加，也更容易有過擬合狀況產生。

Accuracy:

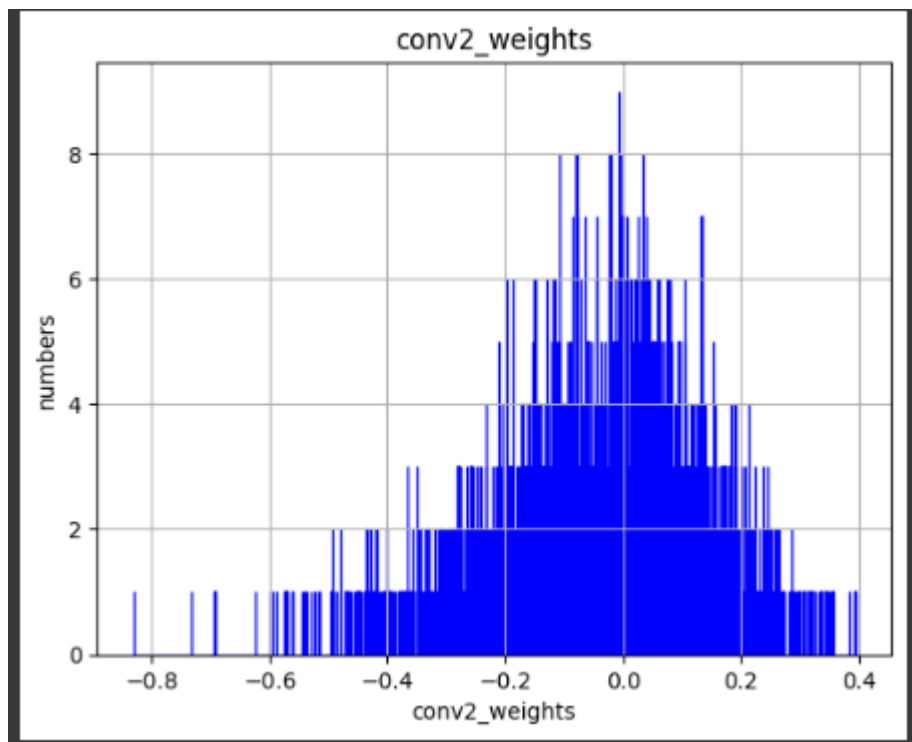
```
[1, 2000] loss: 0.379
[1, 4000] loss: 0.134
[1, 6000] loss: 0.114
[1, 8000] loss: 0.095
[1, 10000] loss: 0.087
[1, 12000] loss: 0.090
[1, 14000] loss: 0.074
98.38
[2, 2000] loss: 0.059
[2, 4000] loss: 0.066
[2, 6000] loss: 0.060
[2, 8000] loss: 0.070
[2, 10000] loss: 0.055
[2, 12000] loss: 0.062
[2, 14000] loss: 0.065
98.67
Finished Training
Accuracy of the network on the test images: 98.67%
```

(2.1.1)

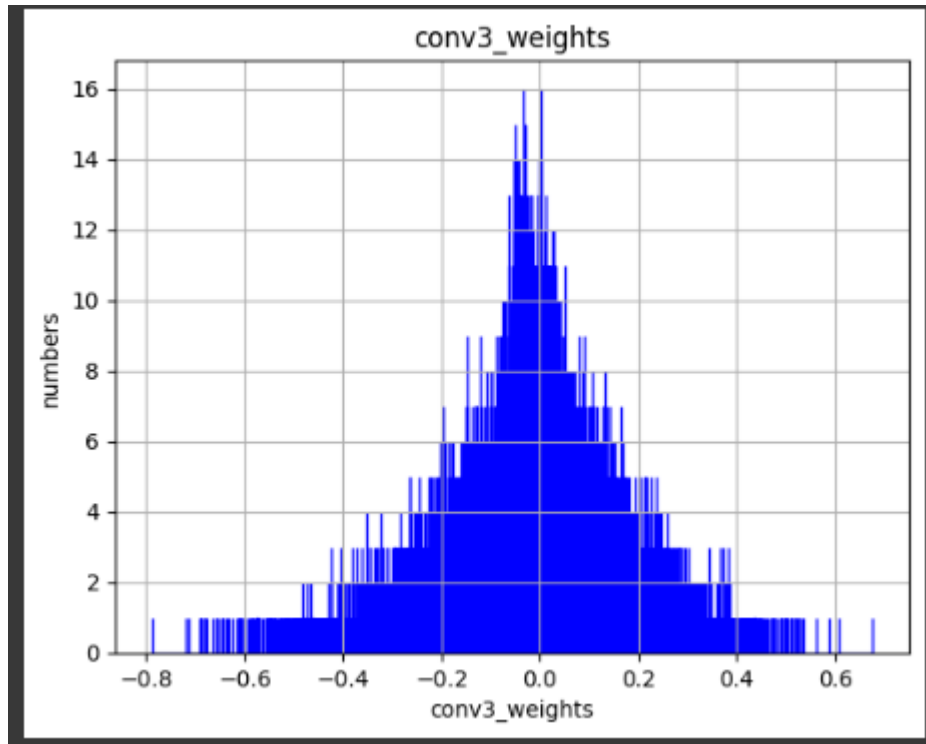
Conv1_weights:




Conv2_weights:



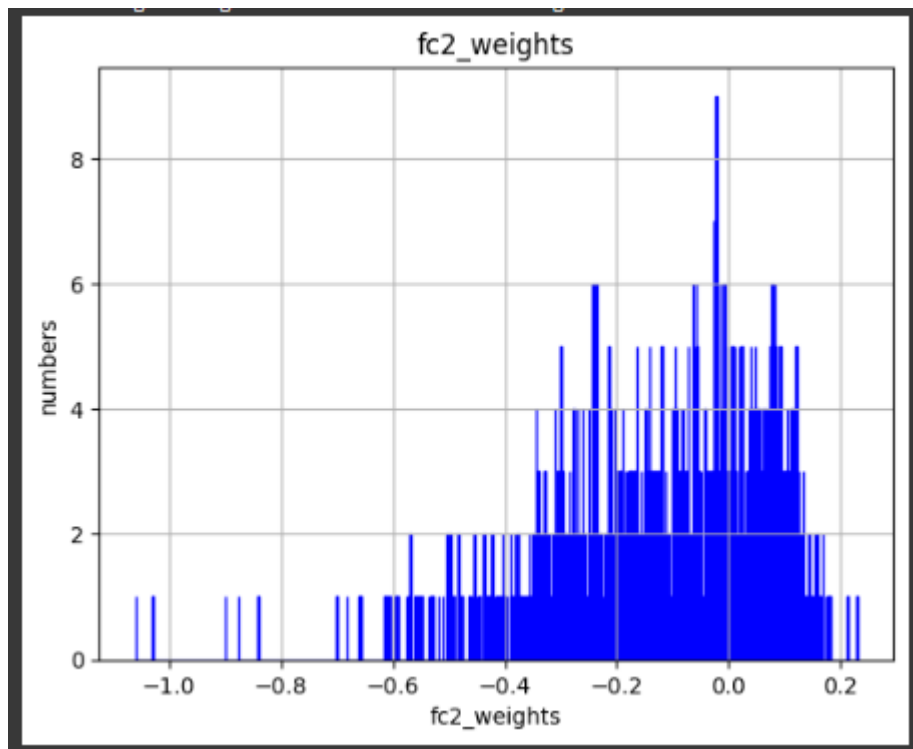
Conv3_weights:



Fc1_weights:



Fc2_weights:



(2.1.2)

Conv1 3-sigma min max range = [-0.4416758120059967, 0.4606795012950897]

Conv1 3-sigma range = [0.9023553133010864]

Conv1 real_value min max range = [-0.35663148760795593, 0.47022220492362976]

Conv1 real_value range = [0.8268536925315857]

Conv1: 3-sigma range is **larger** than real range

Conv2 3-sigma min max range = [-0.5404166579246521, 0.45148125290870667]

Conv2 3-sigma range = [0.9918979406356812]

Conv2 real_value min max range = [-0.8315535187721252, 0.3938743770122528]

Conv2 real_value range = [1.2254278659820557]

Conv2: 3-sigma range is **smaller** than real range

Conv3 3-sigma min max range = [-0.4285754859447479, 0.3955545723438263]

Conv3 3-sigma range = [0.8241300582885742]

Conv3 real_value min max range = [-0.7897406220436096, 0.6771494150161743]

Conv3 real_value range = [1.4668900966644287]

Conv3: 3-sigma range is **smaller** than real range

Fc1 3-sigma min max range = [-0.4010276794433594 , 0.3817105293273926]

Fc1 3-sigma range = [0.782738208770752]

Fc1 real_value min max range = [-0.600527822971344 , 0.47491809725761414]

Fc1 real_value range = [1.0754458904266357]

Fc1: 3-sigma range is **smaller** than real range

Fc2 3-sigma min max range = [-0.6948025822639465 , 0.4443071484565735]

Fc2 3-sigma range = [1.13910973072052]

Fc2 real_value min max range = [-1.0603604316711426 , 0.22938837110996246]

Fc2 real_value range = [1.2897487878799438]

Fc2: 3-sigma range is **smaller** than real range

(2.1.3)

在 3-sigma range 和真實值的 range 比較中，我會選擇真實值的 range，原因在於它的範圍較廣，較能涵蓋全部的值，且能 clamp 掉極值。

(2.2.1)

Sw 是用來量化權重的 scaling factor，在本次題目中，要實作的是 symmetric quantization，因此我們需要找到|Weight|max，接著題目需要我們 fix-point 在 8-bits，帶入下面公式，即可求得 Sw。

$$Sw = \frac{2^n - 1}{2|Weight|max} , n = 8$$

(2.2.2)

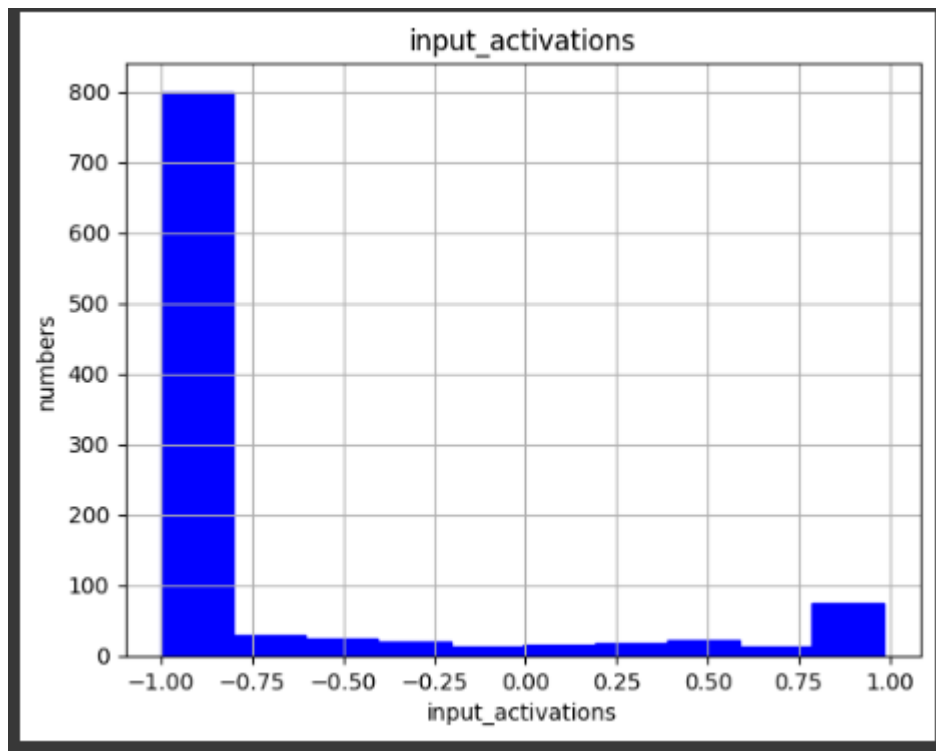
在量化過程中，會從 32-bits floating point 轉成 8-bits fixed point，雖然可以有效減少運算量和資源的消耗，但會犧牲一點準確度，造成準確度下降

Accuracy:

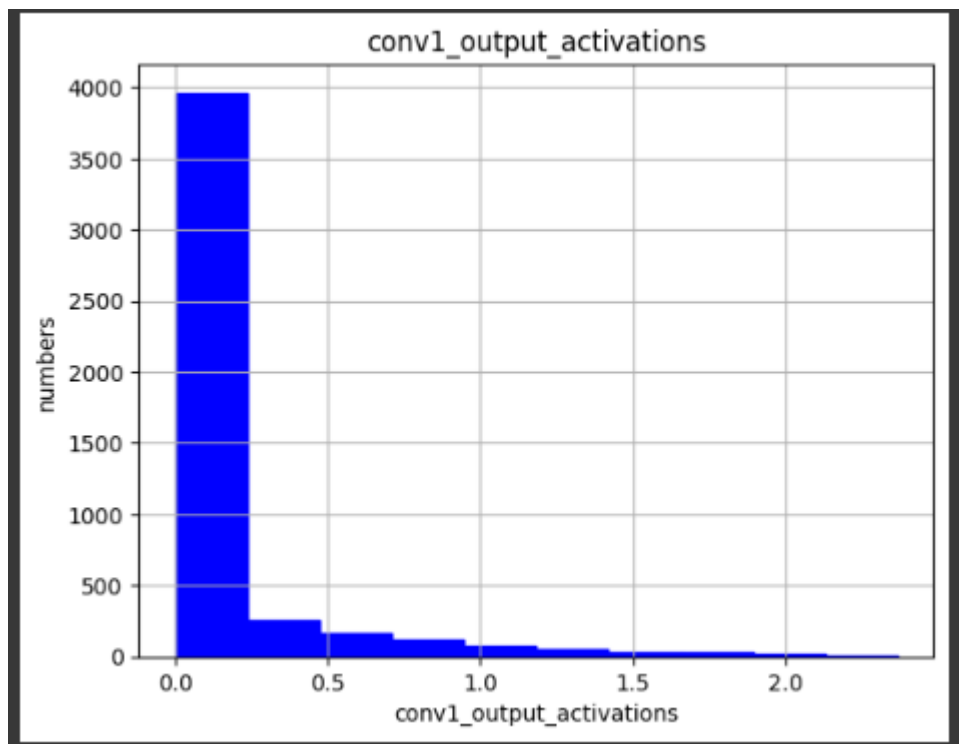
Accuracy of the network after quantizing all weights: 98.67%

(2.3.1)

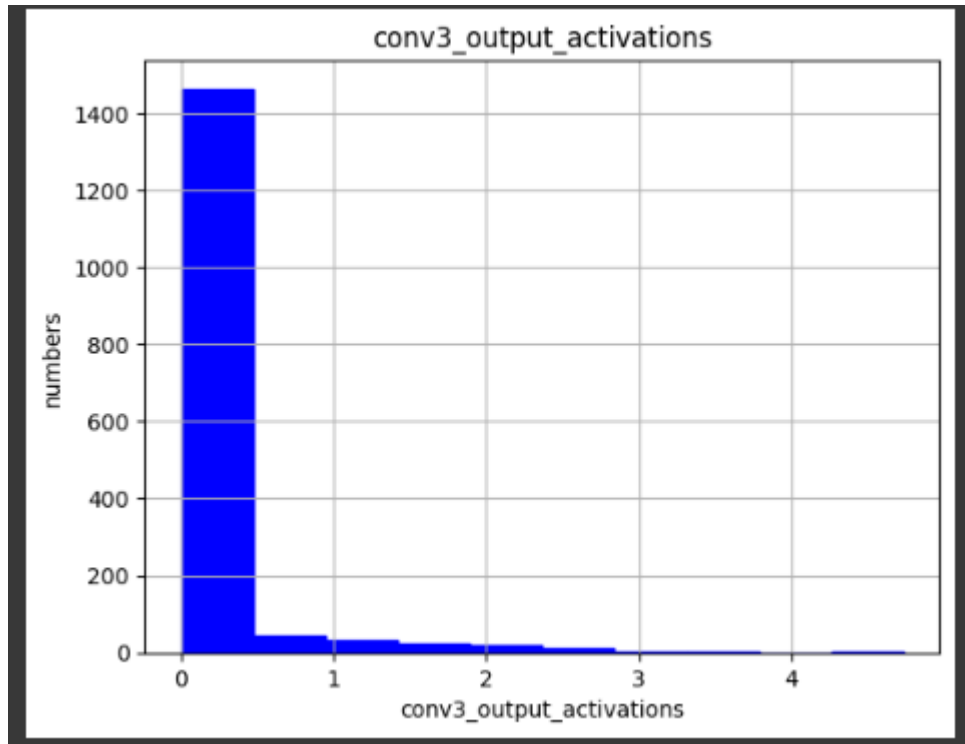
Input activations:



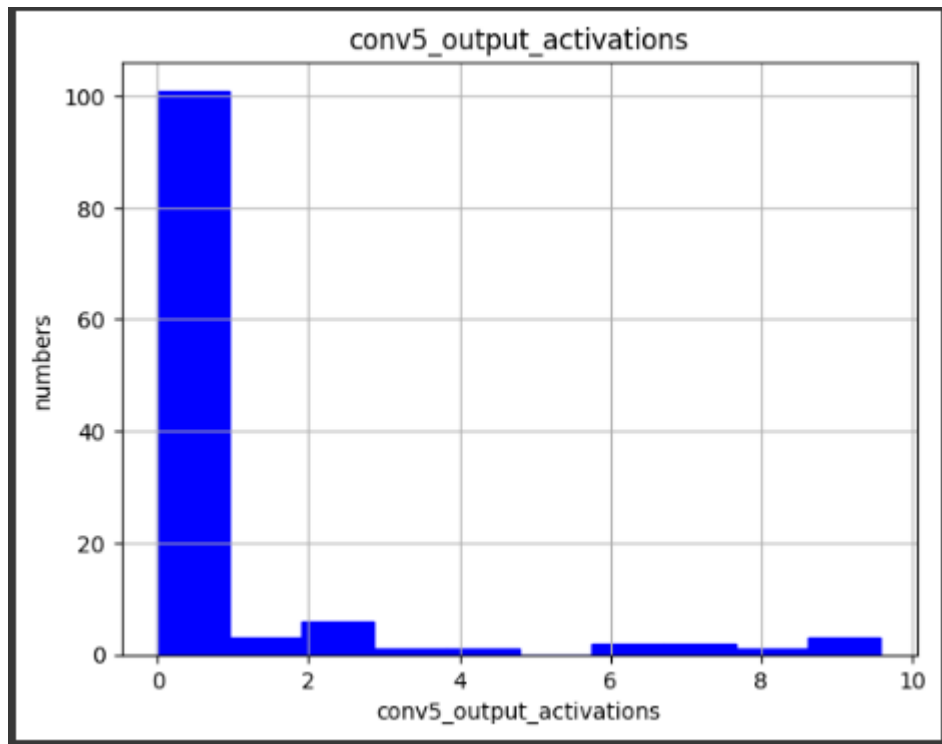
Conv1_output_activations:



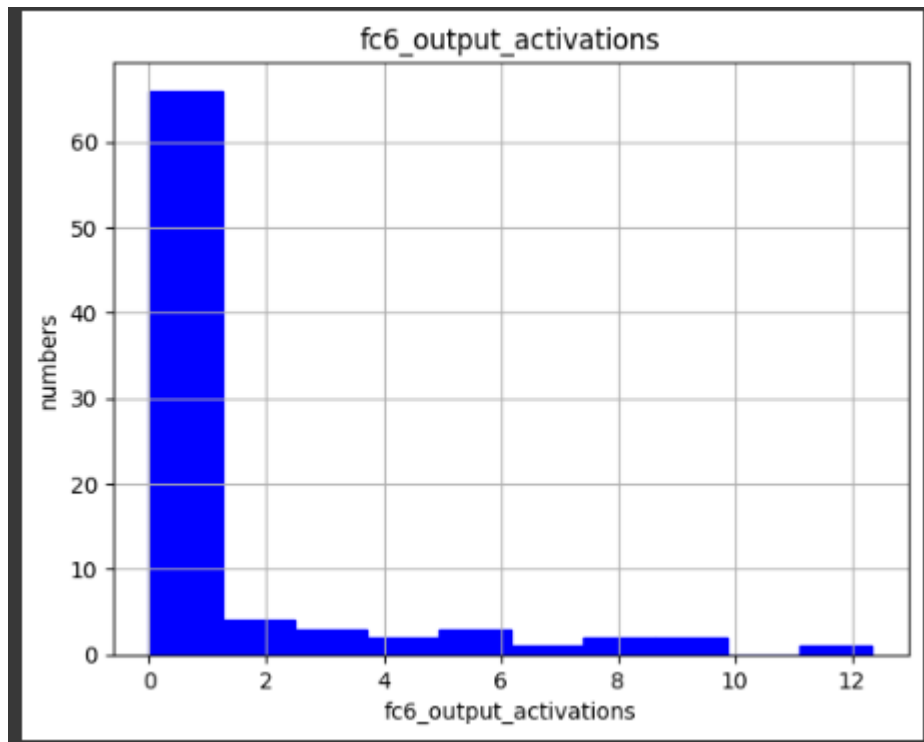
Conv3_output_activations:



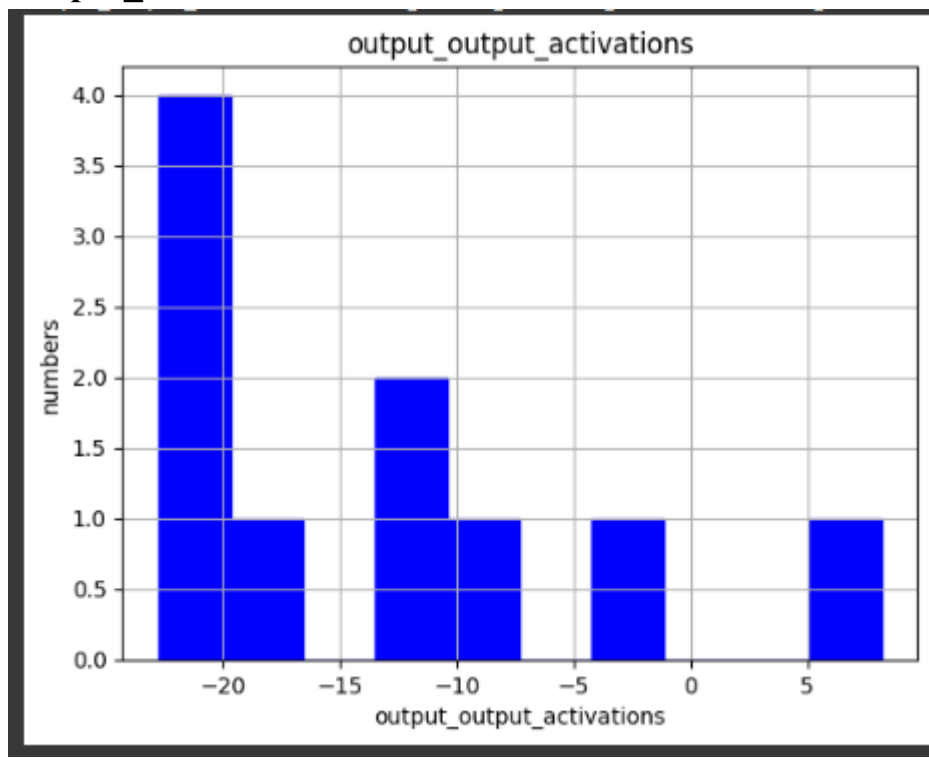
Conv5_output_activations:



Fc6_output_activations:



Output_activations:



(2.3.2)

input_activations 3-sigma range = [-2.484403610229492 , 1.0379424095153809]

input_activations 3-sigma range = [3.522346019744873]

input_activations real_value range = [-1.0 , 0.9843137264251709]

input_activations real_value range = [1.984313726425171]

input_activations: 3-sigma range is **larger** than real range

conv1_output_activations 3-sigma range = [-0.8579810857772827,
1.1403552293777466]

conv1_output_activations 3-sigma range = [1.9983363151550293]

conv1_output_activations real_value range = [0.0 , 2.3684353828430176]

conv1_output_activations real_value range = [2.3684353828430176]

conv1_output_activations: 3-sigma range is **smaller** than real range

conv3_output_activations 3-sigma range = [-1.2368762493133545 ,
1.5076217651367188]

conv3_output_activations 3-sigma range = [2.7444980144500732]

conv3_output_activations real_value range = [0.0 , 4.735324859619141]

conv3_output_activations real_value range = [4.735324859619141]

conv3_output_activations: 3-sigma range is **smaller** than real range

conv5_output_activations 3-sigma range = [-5.406546115875244 ,
6.932685375213623]

conv5_output_activations 3-sigma range = [12.339231491088867]

conv5_output_activations real_value range = [0.0 , 9.57634449005127]

conv5_output_activations real_value range = [9.57634449005127]

conv5_output_activations: 3-sigma range is **larger** than real range

fc6_output_activations 3-sigma range = [-6.647026538848877 , 9.058034896850586]

fc6_output_activations 3-sigma range = [15.705060958862305]

fc6_output_activations real_value range = [0.0 , 12.349260330200195]

fc6_output_activations real_value range = [12.349260330200195]

fc6_output_activations: 3-sigma range is **larger** than real range

output_activations 3-sigma range = [-43.895992279052734 , 17.61233901977539]

output_activations 3-sigma range = [61.508331298828125]

output_activations real_value range = [-22.77204132080078 , 8.168824195861816]

output_activations real_value range = [30.94086456298828]

output_activations: 3-sigma range is **larger** than real range

(2.3.3)

在 3-sigma range 和真實值的 range 比較中，這裡我會選擇 3-sigma range，原因在於它的範圍較廣，較能涵蓋全部的值，且能 clamp 掉極值。

(2.4.1)

求得 S_I , Sw_{conv1} , So_{conv1} 的值是透過課堂上所教的方法，透過找出 Input, $Weight_{conv1}$, $Output_{conv1}$ 絕對值後的最大值，乘上兩倍後，透過想要求得的 bits 數 (這裡是 8-bits)，透過下列方程式得到 scaling_factor。

$$S_I = \frac{255}{2 * |Input|_{max}}$$
$$Sw_{conv1} = \frac{255}{2 * |Weight_{conv1}|_{max}}$$
$$S_I = \frac{255}{2 * |Output_{conv1}|_{max}}$$

(2.4.2)

$$Sw_{conv1} * W_{conv1q} * S_I * I_q = So_{conv1} * O_{conv1q}$$

$$\frac{Sw_{conv1} * S_I}{So_{conv1}} * W_{conv1q} * I_q = O_{conv1q}$$

$$M1 = \frac{Sw_{conv1} * S_I}{So_{conv1}}$$

where W_{conv1q} is the quantized 8-bit signed integer weight tensor,

I_q is the quantized 8-bit signed integer input activation tensor,

and O_{conv1q} is the quantized 8-bit signed integer output activation tensor.

(2.4.3)

$$Sw_{conv3} * W_{conv3} * So_{conv1} * O_{conv1} = So_{conv3} * O_{conv3q}$$

$$M3 = \frac{Sw_{conv3} * So_{conv1}}{So_{conv3}}$$

(2.4.4)

$$M = \frac{Sw(\text{當前層}) * SI(\text{前一層})}{So(\text{當前層})}$$

Accuracy:

```
Accuracy of the network after quantizing both weights and activations: 98.69%
```

(2.4.6)

使用 floor 函式能夠將小數向下取整到最接近的整數，使量化值始終小於或等於原始值，它的好處是保持量化的一致性，確保他們在某個範圍中，也使數值不會分布偏向較大的大小。

(2.4.7)

能夠有效的減少浮點數的運算，讓數值不會過度縮放

(2.5.1)

$$S_w * W_q * S_I * I_q + S_B * B_q = S_o * O_q$$
$$\frac{S_w * S_I}{S_o} * \left(W_q * I_q + \frac{S_b}{S_w * S_I} * B_q \right) = O_q$$
$$M = \frac{S_w * S_I}{S_o}$$

Accuracy:

```
[1, 2000] loss: 0.420
[1, 4000] loss: 0.144
[1, 6000] loss: 0.110
[1, 8000] loss: 0.104
[1, 10000] loss: 0.084
[1, 12000] loss: 0.091
[1, 14000] loss: 0.079
98.23
[2, 2000] loss: 0.071
[2, 4000] loss: 0.056
[2, 6000] loss: 0.067
[2, 8000] loss: 0.073
[2, 10000] loss: 0.057
[2, 12000] loss: 0.058
[2, 14000] loss: 0.063
97.77
Finished Training
Accuracy of the network (with a bias) on the test images: 97.77%
```

```
Accuracy of the network on the test images after all the weights are quantized but the bias isn't: 97.81%
```

```
Accuracy of the network on the test images after all the weights and the bias are quantized: 97.77%
```

(3.1.1)

QAT 能夠達到較高的準確率是因為它在訓練期間加入了量化操作，讓量化能夠更適應訓練的過程，模型在訓練時，就已經可以在低精度上進行運算，因此，模型就可以最大程度的減少準確率的損失，而 PTQ 相反，它是在訓練完成後對模型進行量化，模型是在高精度上訓練，因此，在訓練完成後，會導致一定程度的準確率損失，然而，QAT 通常會比 PTQ 來的更耗時且複雜。

(3.1.2)

Quant layer: 用於將浮點數轉換成 fixed point

Dequant Layer: 量化後的整數值轉為浮點數值

Accuracy:

```
[1, 2000] loss: 0.541
[1, 4000] loss: 0.125
[1, 6000] loss: 0.111
[1, 8000] loss: 0.088
[1, 10000] loss: 0.084
[1, 12000] loss: 0.079
[1, 14000] loss: 0.074
98.16
[2, 2000] loss: 0.058
[2, 4000] loss: 0.058
[2, 6000] loss: 0.067
[2, 8000] loss: 0.063
[2, 10000] loss: 0.055
[2, 12000] loss: 0.067
[2, 14000] loss: 0.065
98.08
Finished Training
Accuracy of the MODEL_FP32: 98.09%
Accuracy of the quantized LeNet-5 model on the test images: 98.11%
```

```
Accuracy of the network with fixed point scale: 98.14%
```